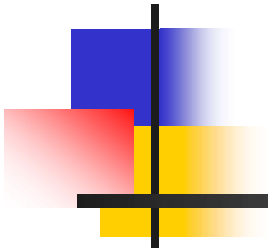



Graphs



Amiya Ranjan Panda

What is a graph?

- A data structure that consists of a set of **nodes** (*vertices*) and a set of edges that relate the nodes to each other.
- The set of edges describes relationships among the vertices.





Formal Definition of Graphs

A graph G is defined as follows: $G=(V, E)$


$V(G)$: a finite, nonempty set of vertices

$E(G)$: a set of edges (pairs of vertices)

a set E that is a subset of $V \times V$ i.e. E is a set of pairs of the form (x, y) where x and y are nodes in V .


- An undirected graph is one in which the pair of vertices in a edge is unordered, $(v_0, v_1) = (v_1, v_0)$
- A directed graph is one in which each edge is a directed pair of vertices, $(v_0, v_1) \neq (v_1, v_0)$

Examples for Graph




G_1

complete graph



G_2

incomplete graphs



G_3

$$V(G_1) = \{v_0, v_1, v_2, v_3\}$$

$$V(G_2) = \{v_0, v_1, v_2, v_3, v_4, v_5, v_6\}$$

$$V(G_3) = \{v_0, v_1, v_2\}$$

$$E(G_1) = \{(v_0, v_1), (v_0, v_2), (v_0, v_3), (v_1, v_2), (v_1, v_3), (v_2, v_3)\}$$

$$E(G_2) = \{(v_0, v_1), (v_0, v_2), (v_1, v_3), (v_1, v_4), (v_2, v_5), (v_2, v_6)\}$$

$$E(G_3) = \{(v_0, v_1), (v_1, v_0), (v_1, v_2)\}$$

complete undirected graph: $n(n-1)/2$ edges

complete directed graph: $n(n-1)$ edges



Complete Graph


- A **complete graph** is a graph that has the maximum number of edges
 - For **undirected graph** with n vertices: maximum number of edges: $n(n-1)/2$
 - For **directed graph** with n vertices, the maximum number of edges: $n(n-1)$
 - **Example**: G1 is a complete graph.




Adjacent and Incident

- If (v_0, v_1) is an edge in an undirected graph,
 - v_0 and v_1 are adjacent
 - The edge (v_0, v_1) is incident on vertices v_0 and v_1
- If (v_0, v_1) is an edge in a directed graph
 - v_0 is adjacent to v_1 , and v_1 is adjacent from v_0 .
 - The edge (v_0, v_1) is incident on v_0 and v_1 .

Example: Graph with feedback loops and a multigraph



(a) self edge



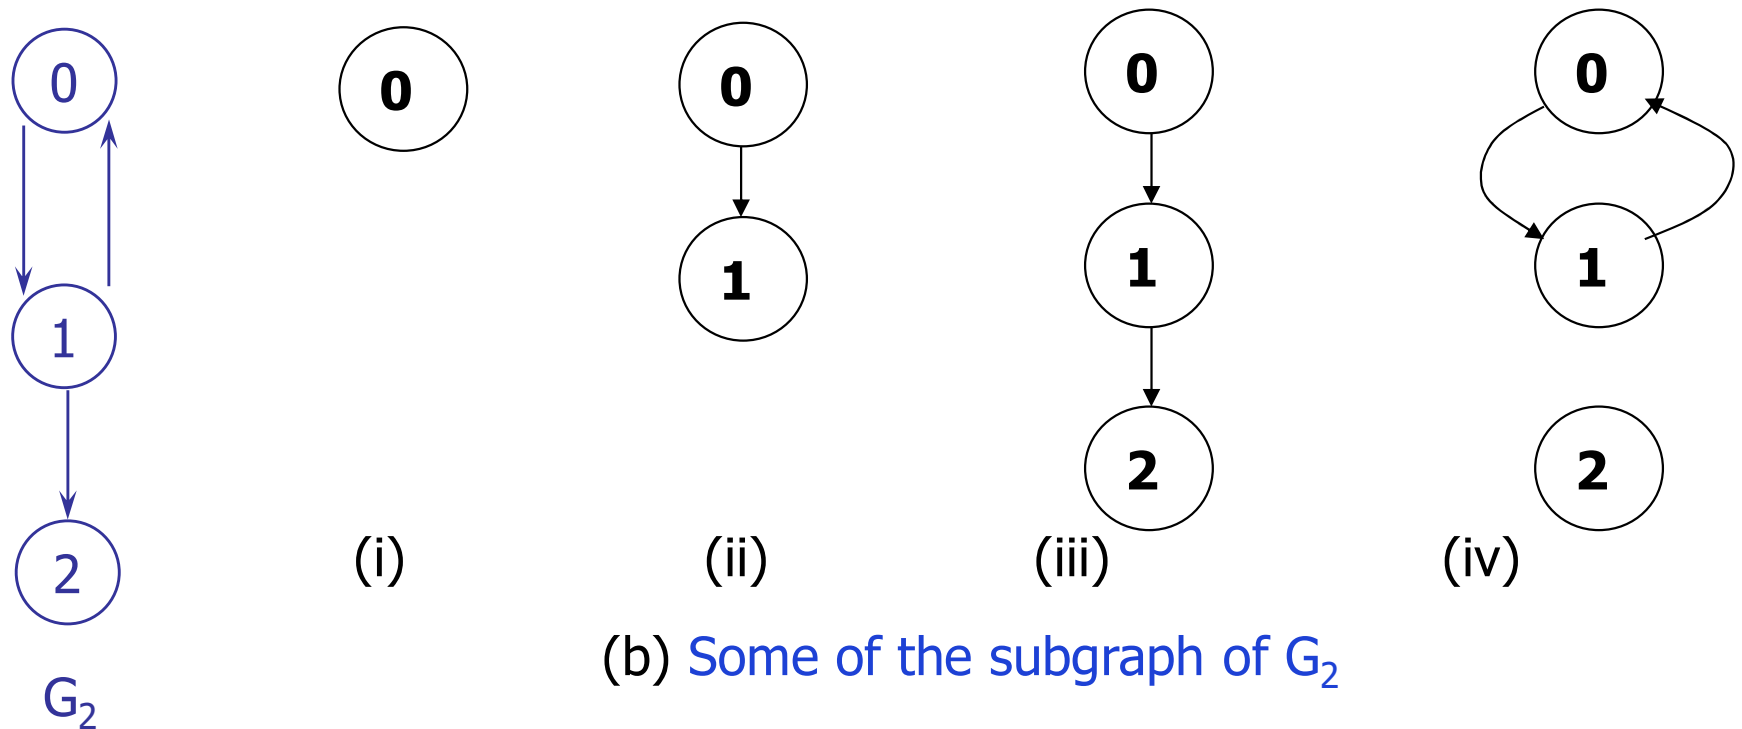
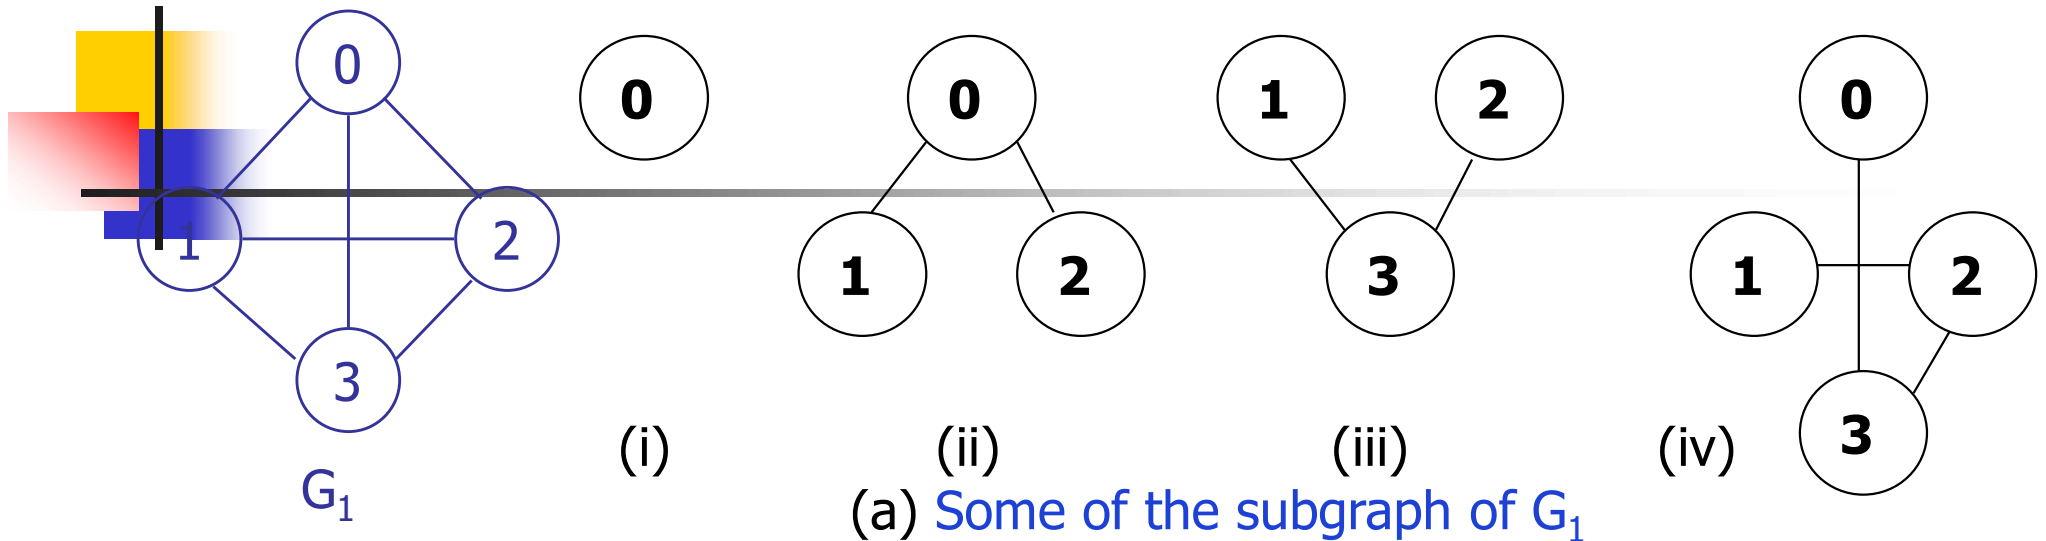
(b) multigraph:
multiple occurrences
of the same edge



Subgraph and Path

- A subgraph of G is a graph, S , such that $V(S)$ is a subset of $V(G)$ and $E(S)$ is a subset of $E(G)$.
- A path from vertex v_p to vertex v_q in a graph G , is a sequence of vertices, $v_p, v_{i1}, v_{i2}, \dots, v_{in}, v_q$ such that $(v_p, v_{i1}), (v_{i1}, v_{i2}), \dots, (v_{in}, v_q)$ are edges in an undirected graph.
- The length of a path is the number of edges on the path.

Subgraphs of G_1 and G_3






Some Terminology


- A simple path is a path in which all vertices, except possibly the first and the last, are distinct.
- A cycle is a simple path in which the first and the last vertices are the same.
- In an undirected graph G , two vertices, v_0 and v_1 , are connected if there is a path from v_0 to v_1 in G .
- An undirected graph is connected if, for every pair of distinct vertices v_i, v_j , there is a path from v_i to v_j .



Connected



G_1



G_2

tree (acyclic graph)




Connected Component

- A connected component or simply component of an undirected graph is a subgraph in which each pair of vertices is connected with each other via a path.
- A tree is a graph that is connected and acyclic.
- A directed graph is strongly connected if there is a directed path from v_i to v_j and also from v_j to v_i .


A graph with two connected components

connected components




Strongly connected components

not strongly connected



strongly connected components







Degree

The degree of a vertex is the number of edges incident to that vertex

- For directed graph,
 - the in-degree of a vertex v is the number of edges that have v as the head
 - the out-degree of a vertex v is the number of edges that have v as the tail
 - if d_i is the degree of a vertex i in a graph G with n vertices and e edges, the number of edges is


$$e = \left(\sum_{i=0}^{n-1} d_i \right) / 2$$

Degrees in Graph



directed graph

in-degree
out-degree





ADT for Graph

structure of Graph

objects: a nonempty set of vertices and a set of undirected edges, where each edge is a pair of vertices

functions: for all graph \in Graph, v , v_1 and $v_2 \in$ Vertices

Graph create(): return an empty graph

Graph insertVertex(graph, v): return a graph with v inserted. v has no incident edge.

Graph insertEdge(graph, v_1 , v_2): return a graph with new edge between v_1 and v_2

Graph deleteVertex(graph, v): return a graph in which v and all edges incident to it are removed

Graph deleteEdge(graph, v_1 , v_2): return a graph in which the edge (v_1 , v_2) is removed

Boolean isEmpty(graph): if (graph==empty graph) return TRUE
else return FALSE

List adjacent(graph, v): return a list of all vertices that are adjacent to v



Graph Representations


- Adjacency Matrix
- Adjacency Lists
- Adjacency Multilists



Adjacency Matrix


- Let $G=(V, E)$ be a graph with n vertices.
- The adjacency matrix of G is a two-dimensional $n \times n$ array, say adjMat .
- If the edge (v_i, v_j) is in $E(G)$, $\text{adjMat}[i][j]=1$
- If there is no such edge in $E(G)$, $\text{adjMat}[i][j]=0$
- The adjacency matrix for an undirected graph is symmetric; the adjacency matrix for a digraph need not be symmetric

Examples for Adjacency Matrix





$$\begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$

G1



$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

G2



$$\begin{bmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$

G4

symmetric



Merits of Adjacency Matrix

- From the adjacency matrix, to determine the connection of vertices is easy.


- The degree of a vertex is $\sum_{j=0}^{n-1} adjMat[i][j]$

- For a **digraph**, the row sum is the out_degree, while the column sum is the in_degree

$$ind(vi) = \sum_{j=0}^{n-1} A[j, i]$$

$$outd(vi) = \sum_{j=0}^{n-1} A[i, j]$$


Adjacency Matrix Properties



| | 1 | 2 | 3 | 4 | 5 |
|---|---|---|---|---|---|
| 1 | 0 | 1 | 0 | 1 | 0 |
| 2 | 1 | 0 | 0 | 0 | 1 |
| 3 | 0 | 0 | 0 | 0 | 1 |
| 4 | 1 | 0 | 0 | 0 | 1 |
| 5 | 0 | 1 | 1 | 1 | 0 |

- Diagonal entries are zero.
- Adjacency matrix of an undirected graph is symmetric.
 - $A(i, j) = A(j, i)$ for all i and j .

Adjacency Matrix (Digraph)



| | 1 | 2 | 3 | 4 | 5 |
|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 0 |
| 2 | 1 | 0 | 0 | 0 | 1 |
| 3 | 0 | 0 | 0 | 0 | 0 |
| 4 | 0 | 0 | 0 | 0 | 1 |
| 5 | 0 | 1 | 1 | 0 | 0 |

- Diagonal entries are zero.
- Adjacency matrix of a digraph need not be symmetric.







Data Structures for Adjacency Lists

Each row in adjacency matrix is represented as an adjacency list.


```
#define MAX_VERTICES 50
```

```
struct node {  
    int vertex;  
    struct node *link;  
};
```


```
struct node * graph[MAX_VERTICES];  
int n=0;           // vertices currently in use
```

G1



G3



G2



Array Length = **n**

of chain nodes = **2e** (undirected graph)

of chain nodes = **e** (digraph)

Array Adjacency Lists

- Each adjacency list is an array list.



Array Length = n

of list elements = $2e$ (undirected graph)

of list elements = e (digraph)



Weighted Graphs

- Cost adjacency matrix.
 - $C(i, j)$ = cost of edge (i, j)
- Adjacency lists => each list element is a pair (adjacent vertex, edge weight)



Graph Traversal Techniques

- The connectivity problem, as well as many other graph problems, can be solved using graph traversal techniques
- There are two standard graph traversal techniques:
 - Depth-First Search (DFS)
 - Breadth-First Search (BFS)



Graph Traversal (Contd.)

- In both DFS and BFS, the nodes of the undirected graph are visited in a systematic manner so that every node is visited exactly once.
- Both BFS and DFS give rise to a tree:
 - When a node x is visited, it is labeled as visited, and it is added to the tree
 - If the traversal got to node x from node y , y is viewed as the parent of x , and x a child of y .



Depth-First Search

DFS follows the following rules:


1. Select an unvisited node x , visit it, and treat as the **current node**
2. Find an unvisited neighbor of the current node, visit it, and make it the new current node;
3. If the current node has no unvisited neighbors, **backtrack** to the its parent, and make that parent the new current node;
4. Repeat steps 2 and 3 until no more nodes can be visited.
5. If there are still unvisited nodes, repeat from step 1.



DFS (Pseudo Code)


```
DFS(input: Graph G, Node v) {  
    if (v == NULL)  
        return;  
    push(v);  
    while (stack is not empty) {  
        pop(v);  
        if (v has not yet been visited)  
            mark_and_visit(v);  
        for (each w adjacent to v)  
            if (w has not yet been visited)  
                push(w);  
    }  
}
```

Example





Policy: Visit adjacent nodes in increasing index order

DFS: Start with Node 5




5 1 0 3 2 7 6 4

DFS: Start with Node 5



push (5)


DFS: Start with Node 5



pop/visit/mark (5)

5



DFS: Start with Node 5



push (2), push (1)

5


DFS: Start with Node 5



pop/visit/mark (1)

5 1

DFS: Start with Node 5





| |
|---|
| 0 |
| 2 |
| 4 |
| 2 |

pop/visit/mark (1)

5 1



DFS: Start with Node 5



pop/visit/mark(0)

5 1 0



DFS: Start with Node 5



push(7), push(3)

5 1 0


DFS: Start with Node 5



pop/visit/mark(3)

5 1 0 3

DFS: Start with Node 5




| |
|---|
| 2 |
| 7 |
| 2 |
| 4 |
| 2 |

push(2)

5 1 0 3

DFS: Start with Node 5





| |
|---|
| 7 |
| 2 |
| 4 |
| 2 |

pop/mark/visit(2)

5 1 0 3 2



DFS: Start with Node 5



pop/mark/visit(7)

5 1 0 3 2 7



DFS: Start with Node 5



push(6)

5 1 0 3 2 7



DFS: Start with Node 5



pop/mark/visit(6)

5 1 0 3 2 7 6


DFS: Start with Node 5



pop/do not visit(2)

5 1 0 3 2 7 6



DFS: Start with Node 5



pop/mark/visit(6)

5 1 0 3 2 7 6 4


DFS: Start with Node 5




pop/do not visit(2)

5 1 0 3 2 7 6 4

DFS: Start with Node 5




5 1 0 3 2 7 6 4



Stack empty,
process completed

DFS: Start with Node 5


Note: edge (0, 3) removed



5 1 0 7 6 2 4 3


DFS (Pseudo Code)

Policy: Don't push nodes twice




```
DFS(input: Graph G, Node v) {  
    if (v == NULL)  
        return;  
    push(v);  
    while (stack is not empty) {  
        pop(v);  
        if (v has not yet been visited)  
            mark_and_visit(v);  
        for (each w adjacent to v)  
            if (w has not yet been visited && not yet stacked)  
                push(w);  
    }  
}
```

DFS (Don't push nodes twice).
Start with Node 5



5 1 0 3 7 6 4 2

DFS: Start with Node 5



2 3 6 7 0 4 1 5



Breadth-First Search

BFS follows the following rules:


1. Select an unvisited node x , visit it, have it be the root in a BFS tree being formed. Its level is called the current level.
2. From each node z in the current level, in the order in which the level nodes were visited, visit all the unvisited neighbors of z .
3. The newly visited nodes from this level form a new level that becomes the next current level.
4. Repeat step 2 & 3 until no more nodes can be visited.
5. If there are still unvisited nodes, repeat from Step 1.



BFS (Pseudo Code)


```
BFS(input: graph G, Node v) {  
    Queue Q;    Integer x, z, y;  
    if (v == NULL)    return;  
    enqueue(v);  
    while (queue is not empty) {  
        dequeue(v);  
        if (v has not yet been visited)  
            mark_and_visit(v);  
        for (each w adjacent to v)  
            if (w has not yet been visited && has not been queued)  
                enqueue(w);  
    }  
}
```


BFS: Start with Node 5




5 1 2 0 4 3 7 6

BFS: Start with Node 5




5

BFS: Node one-away




BFS: Visit 1 and 2




5 1 2

BFS: Nodes two-away




BFS: Visit 0 and 4




5 1 2 0 4

BFS: Nodes three-away




5 1 2 0 4

BFS: Visit nodes 3 and 7




5 1 2 0 4 3 7

BFS: Node four-away



5 1 2 0 4 3 7


BFS: Visit 6



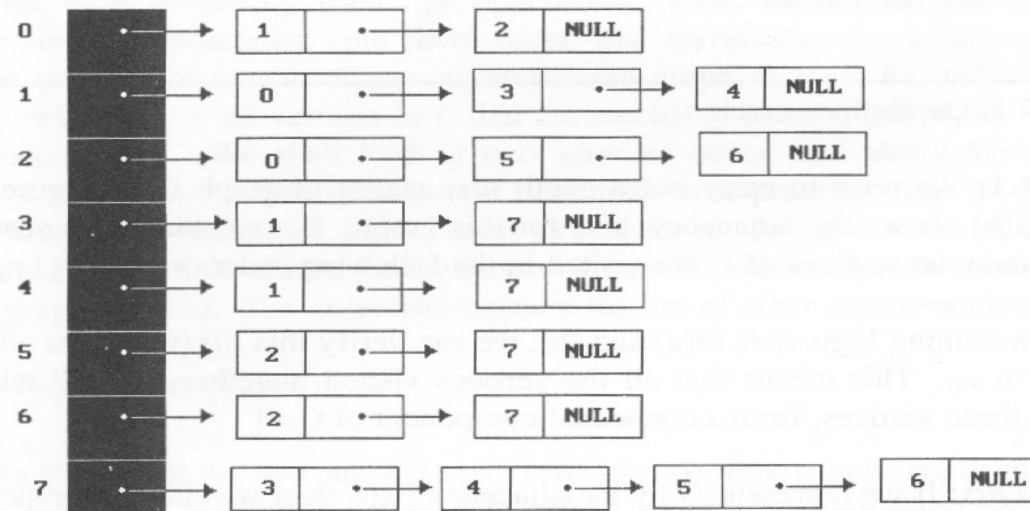
5 1 2 0 4 3 7 6

Graph G and its adjacency lists

depth first search: v0, v1, v3, v7, v4, v5, v2, v6



(a)



(b)

breadth first search: v0, v1, v2, v3, v4, v5, v6, v7



Topological sort

- Given a **set of tasks** and a **set of dependencies** (**precedence constraints**) of form "task A must be done before task B"
- **Topological sort**: An ordering of the tasks that conforms with the given dependencies
- **Goal**: Find a topological sort of the tasks or decide that there is no such ordering




Topological Sort

- Sorting technique over DAGs (**Directed Acyclic Graphs**)
- It creates a **linear sequence** (**ordering**) for the nodes such that:
 - If u has an outgoing edge to v \rightarrow then u must finish before v starts
- Very common in ordering jobs or tasks

Examples

- **Scheduling:** When scheduling *task graphs* in distributed systems, usually we first need to sort the tasks topologically ...and then assign them to resources.



Topological Sort Example

A job consists of 10 tasks with the following precedence rules:

Must start with 7, 5, 4 or 9.


Task 1 must follow 7.

Tasks 3 & 6 must follow both 7 & 5.


8 must follow 6 & 4.

2 must follow 4.

10 must follow 2.




Topological Sort Example



Topological Sort is not unique

- Topological sort is not unique.
- The following are all topological sort of the graph below:




$s1 = \{a, b, c, d, e, f, g, h, i\}$

$s2 = \{a, c, b, f, e, d, h, g, i\}$

$s3 = \{a, b, d, c, e, g, f, h, i\}$

$s4 = \{a, c, f, b, e, h, d, g, i\}$
etc.

Topological Sort Algorithm



- One way to find a topological sort is to consider in-degrees of the vertices.
- The first vertex must have in-degree zero -- every DAG must have at least one vertex with in-degree zero.
- The Topological sort algorithm is:

```
int topologicalOrderTraversal( ) {  
    int numVisitedVertices = 0;  
    while(there are more vertices to be visited){  
        if(there is no vertex with in-degree 0)  
            break;  
        else {  
            select a vertex v that has in-degree 0;  
            visit v;  
            numVisitedVertices++;  
            delete v and all its emanating edges;  
        }  
    }  
    return numVisitedVertices;  
}
```

Topological Sort Example

- Demonstrating Topological Sort.

